极客时间 训练营

MapReduce与Yarn

张语

目录

重点内容回顾

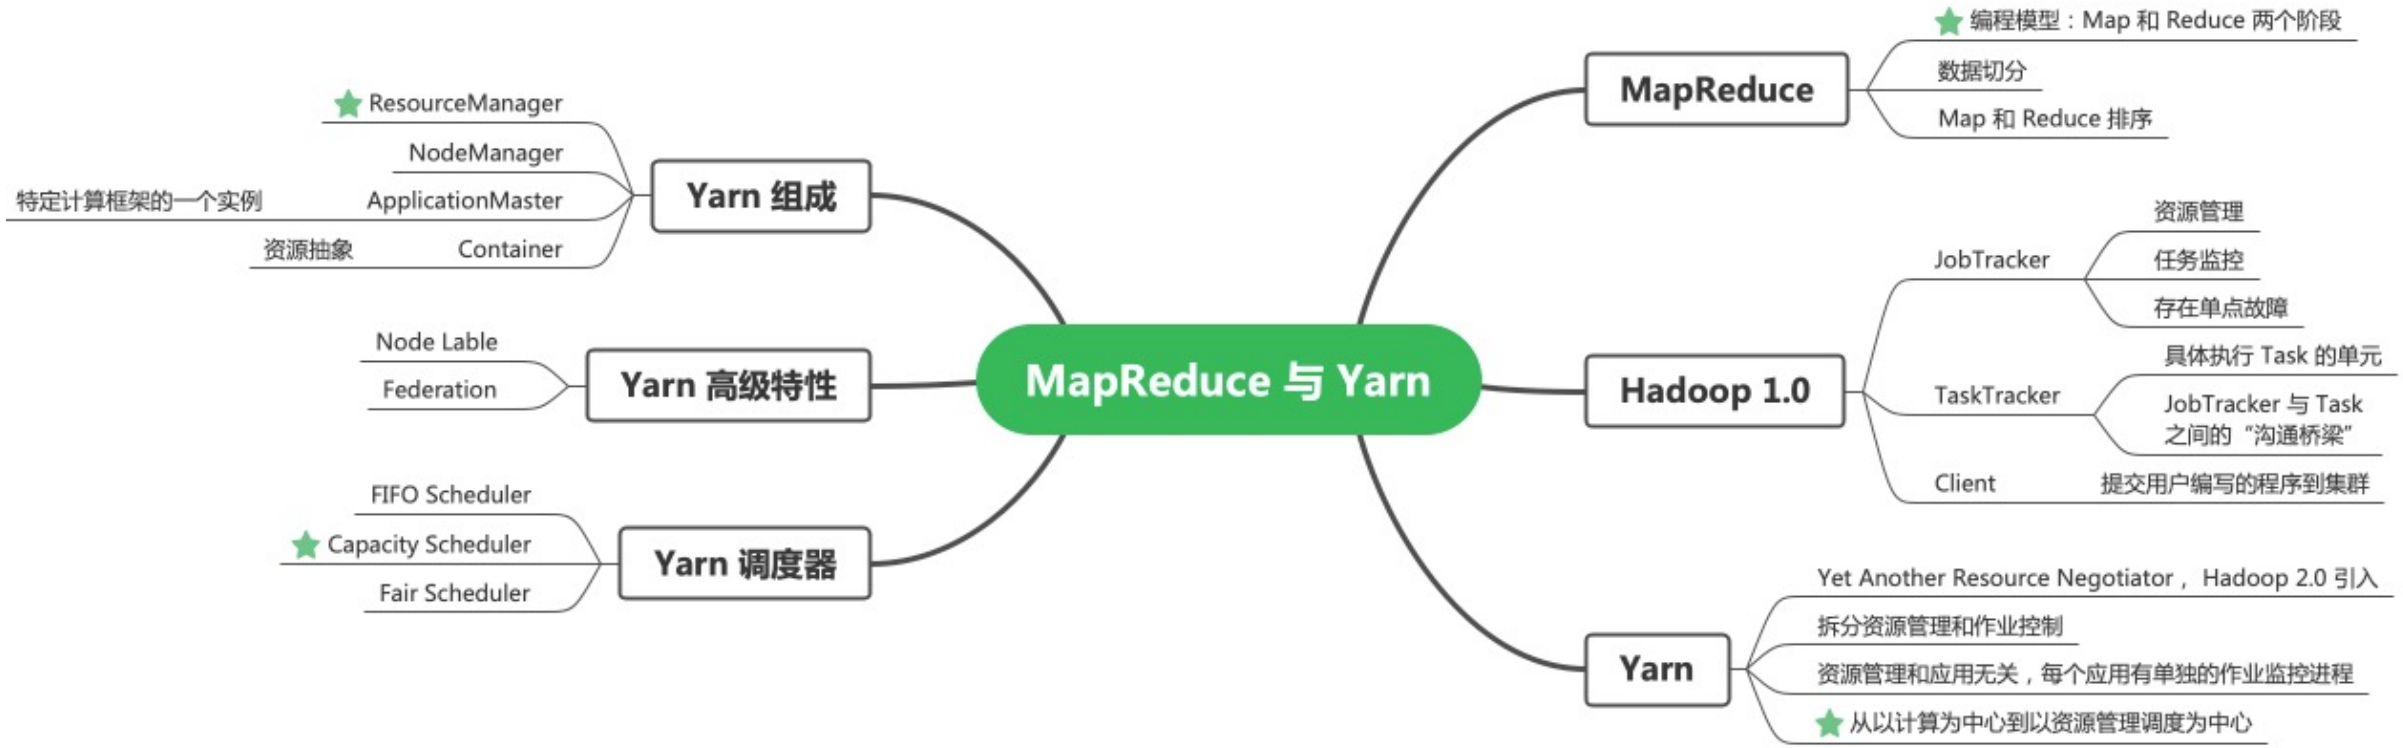
MapReduce 与 Yarn
MapReduce
编程模型：Map 和 Reduce 两个阶段
数据切分
Map 和 Reduce 排序
Hadoop 1.0
JobTracker
资源管理
任务监控
存在单点故障
TaskTracker
具体执行 Task 的单元
JobTracker 与 Task 之间的“沟通桥梁”
Client
提交用户编写的程序到集群
Yarn
Yet Another Resource Negotiator， Hadoop 2.0 引入
拆分资源管理和作业控制
资源管理和应用无关，每个应用有单独的作业监控进程
从以计算为中心到以资源管理调度为中心
Yarn 组成
ResourceManager
NodeManager
ApplicationMaster
特定计算框架的一个实例
Container
资源抽象
Yarn 高级特性
Node Lable
Federation
Yarn 调度器
FIFO Scheduler
Capacity Scheduler
Fair Scheduler

MapReduce编程模型

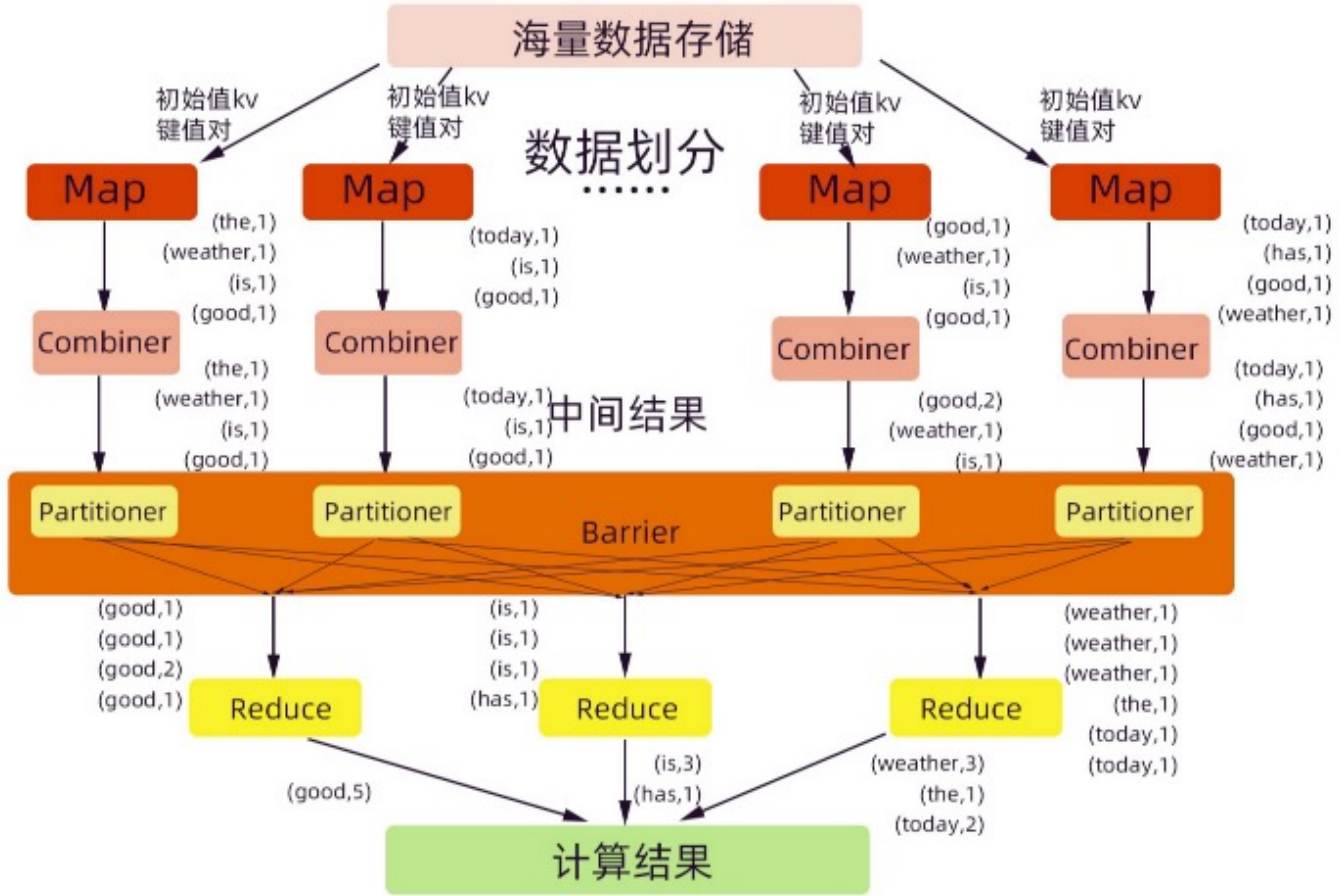

MapReduce编程模型

极客时间训练营

- 分治法
  1. 分解原问题。将原问题分解为若干个规模较小，相互独立，且与原问题形式相同的子问题
  2. 求解子问题。若子问题规模较小且容易被解决则直接求解，否则递归地求解各个子问题
  3. 合并解，就是将各个子问题的解合并为原问题的解
- MapReduce
  - Map：分解，把复杂的任务分解为若干个简单的任务执行
  - Reduce：合并，对 Map 阶段的结果进行汇总

MapReduce编程模型

极客时间训练营

- 使用场景：任务可被分解成相互独立的子问题
  1. 迭代。遍历输入数据，并解析成 key/value 对
  2. 将输入 key/value 对映射（map）成另外一些 key/value 对
  3. 根据 key 对中间数据进行分组（group）
  4. 以组为单位对数据进行归约（reduce）
  5. 迭代。将 reduce 阶段产生的key/value对保存到输出文件中

MapReduce 缺点

极客时间训练营

- **高昂的维护成本**
  - MapReduce 模型过于抽象，当遇到复杂度较高的需求时需要手动编排多个Map和多个Reduce任务
  - Spark、Flink等使用有向无环图（DAG）表达数据处理流程
    - 每个节点都表达一种通用的数据集
    - 每一条边表达一种通用的数据变换
- **时间性能较差**
  - Map 和 Reduce 的输出都需要落盘（本地磁盘、HDFS）

作业讲解

极客时间训练营

- 统计每一个手机号耗费的总上行流量、下行流量、总流量
- 作业详见：https://u.geekbang.org/lesson/180?article=491889
- 基本思路：
  - Map 阶段：
    - 读取一行数据，切分字段。
    - 抽取手机号、上行流量、下行流量。
    - 以手机号为 key，bean 对象为 value 输出，即 context.write(手机号，bean)。
  - Reduce 阶段：
    - 累加上行流量和下行流量得到总流量。
    - 实现自定义的 bean 来封装流量信息，并将 bean 作为 map 输出的 key 来传输

作业讲解

极客时间训练营

hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/Mapper.java

```
public class Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {

  /** The <code>Context</code> passed on to the {@link Mapper} implementations. */
  public abstract class Context
    implements MapContext<KEYIN,VALUEIN,KEYOUT,VALUEOUT> {
  }

  /** Called once at the beginning of the task. */
  protected void setup(Context context
                       ) throws IOException, InterruptedException {}

  /**
   * Called once for each key/value pair in the input split. Most applications
   * should override this, but the default is the identity function.
   */
  @unchecked
  protected void map(KEYIN key, VALUEIN value,
                     Context context) throws IOException, InterruptedException {
    context.write((KEYOUT) key, (VALUEOUT) value);
  }
```

作业讲解

极客时间训练营

hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java/org/apache/hadoop/mapreduce/Reducer.java

```
public class Reducer<KEYIN,VALUEIN,KEYOUT,VALUEOUT> {

  /** The <code>Context</code> passed on to the {@link Reducer} implementations. */
  public abstract class Context
    implements ReduceContext<KEYIN,VALUEIN,KEYOUT,VALUEOUT> {
  }

  /** Called once at the start of the task. */
  protected void setup(Context context
                       ) throws IOException, InterruptedException {}

  /**
   * This method is called once for each key. Most applications will define
   * their reduce class by overriding this method. The default implementation
   * is an identity function.
   */
  /unchecked/
  protected void reduce(KEYIN key, Iterable<VALUEIN> values, Context context
                        ) throws IOException, InterruptedException {
    for(VALUEIN value: values) {
      context.write((KEYOUT) key, (VALUEOUT) value);
    }
  }
```

作业讲解

极客时间训练营

```java
public class PhoneTraffic implements Writable {
    // 上行流量
    private long up;
    // 下行流量
    private long down;
    // 总流量
    private long sum;

    public PhoneTraffic() {}

    public PhoneTraffic(long up, long down, long sum) {...}

    @Override
    public void write(DataOutput dataOutput) throws IOException {
        dataOutput.writeLong(up);
        dataOutput.writeLong(down);
        dataOutput.writeLong(sum);
    }

    @Override
    public void readFields(DataInput dataInput) throws IOException {
        this.up = dataInput.readLong();
        this.down = dataInput.readLong();
        this.sum = dataInput.readLong();
    }
```

作业讲解

极客时间训练营

```
public static class TrafficMapper extends Mapper<Object, Text, Text, PhoneTraffic> {

    public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
        String[] lines = value.toString().split( regex: "\t");
        if (lines.length < 10) {
            return;
        }

        String phone = lines[1];
        try {
            long up = Long.parseLong(lines[8]);
            long down = Long.parseLong(lines[9]);
            context.write(new Text(phone), new PhoneTraffic(up, down, sum: up + down));
        } catch (NumberFormatException e) {
            System.err.println("parseLong failed" + e.getMessage());
        }
    }
}
```

作业讲解

极客时间训练营

```
public static class TrafficReducer extends Reducer<Text,PhoneTraffic,Text,PhoneTraffic> {

    public void reduce(Text key, Iterable<PhoneTraffic> values, Context context
    ) throws IOException, InterruptedException {
        int totalUp = 0;
        int totalDown = 0;
        int sumTraffic = 0;
        for (PhoneTraffic val : values) {
            totalUp += val.getUp();
            totalDown += val.getDown();
            sumTraffic += val.getSum();
        }

        context.write(key, new PhoneTraffic(totalUp, totalDown, sumTraffic));
    }
}
```

作业讲解

极客时间训练营

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();
    if (otherArgs.length < 2) {
        System.err.println("Usage: TrafficStat <in> <out>");
        System.exit( status: 2);
    }
    System.out.println("otherArgs: " + Arrays.toString(otherArgs));

    Job job = Job.getInstance(conf, jobName: "TrafficStat");
    job.setJarByClass(TrafficStat.class);
    job.setMapperClass(TrafficMapper.class);
    job.setCombinerClass(TrafficReducer.class);
    job.setReducerClass(TrafficReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(PhoneTraffic.class);

    job.setNumReduceTasks(1);

    FileInputFormat.addInputPath(job, new Path(otherArgs[otherArgs.length - 2]));
    FileOutputFormat.setOutputPath(job, new Path(otherArgs[otherArgs.length - 1]));
    System.exit(job.waitForCompletion( verbose: true) ? 0 : 1);
}
```

作业讲解

极客时间训练营

运行演示

hadoop jar bigdata-tutorial-1.0-SNAPSHOT.jar TrafficStat /user/public/week2/ /user/public/week2_out/

Hadoop Streaming

极客时间训练营

- 允许用户使用**可执行文件或者脚本**作为 Mapper/Reducer 创建和运行 MapReduce 任务，降低开发门槛
- 要求编写的 Mapper/Reducer 从标准输入中读取数据，将结果写到标准输出中
- 官方文档：https://hadoop.apache.org/docs/stable/hadoop-streaming/HadoopStreaming.html

```
mapred streaming \
        -input myInputDirs \
        -output myOutputDir \
        -mapper /bin/cat \
        -reducer /usr/bin/wc
```

Hadoop Streaming – 配置

极客时间训练营

Parameter	Optional/Required	Description
-input directoryname or filename	Required	Input location for mapper
-output directoryname	Required	Output location for reducer
-mapper executable or JavaClassName	Optional	Mapper executable. If not specified, IdentityMapper is used as the default
-reducer executable or JavaClassName	Optional	Reducer executable. If not specified, IdentityReducer is used as the default
-file filename	Optional	Make the mapper, reducer, or combiner executable available locally on the compute nodes
-combiner streamingCommand or JavaClassName	Optional	Combiner executable for map output
-cmdenv name=value	Optional	Pass environment variable to streaming commands
-numReduceTasks	Optional	Specify the number of reducers

Hadoop Streaming – 原理

- 使用可执行文件或者脚本文件充当Mapper或者Reducer时，Java端的Mapper或者Reducer充当了wrapper角色，将输入文件中的key和value直接传递给可执行文件或者脚本文件进行处理，并将处理结果写入HDFS

Map 阶段

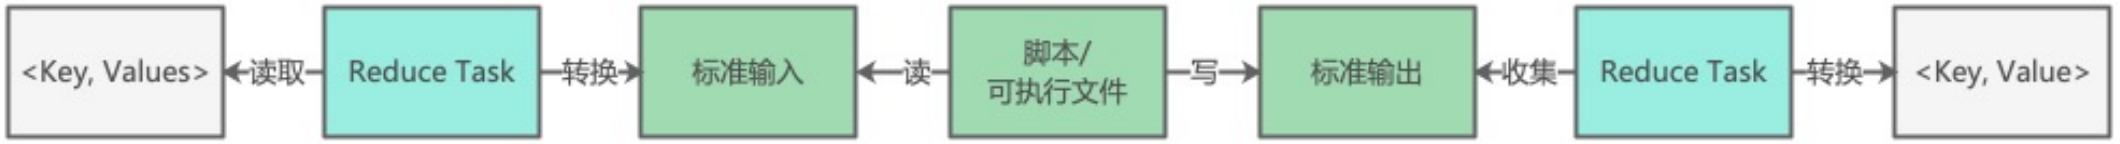

Reduce 阶段

Hadoop Streaming – 原理

难点：如何使用标准输入输出实现Java与可执行文件、脚本之间的通信？

```
205         // Start the process
206         ProcessBuilder builder = new ProcessBuilder(argvSplit);
207         builder.environment().putAll(childEnv.toMap());
208         sim = builder.start();
209
210         clientOut_ = new DataOutputStream(new BufferedOutputStream(
211                                                 sim.getOutputStream(),
212                                                 BUFFER_SIZE));
213         clientIn_ = new DataInputStream(new BufferedInputStream(
214                                                 sim.getInputStream(),
215                                                 BUFFER_SIZE));
216         clientErr_ = new DataInputStream(new BufferedInputStream(sim.getErrorStream()));
```

hadoop-tools/hadoop-streaming/src/main/java/org/apache/hadoop/streaming/PipeMapRed.java

Hadoop Streaming – 使用

极客时间训练营

```
def traffic_mapper():
    for line in sys.stdin:
        arr = line.strip().split("\t")
        if len(arr) < 10:
            continue
        # 电话号码，上行流量，下行流量
        print("%s\t%s\t%s" % (arr[1], arr[8], arr[9]))
```

Hadoop Streaming – 使用

极客时间训练营

```
def traffic_reducer():
    last_phone = ""
    up, down = 0, 0
    for line in sys.stdin:
        arr = line.strip().split("\t")

        if last_phone == "":
            last_phone = arr[0]
        elif last_phone != arr[0]:
            # 输出上一个电话的流量统计
            print("%s\t%s\t%s\t%s" % (last_phone, up, down, up + down))
            last_phone = arr[0]
            up, down = 0, 0

        up += int(arr[1])
        down += int(arr[2])

    if last_phone != "":
        print("%s\t%s\t%s\t%s" % (last_phone, up, down, up + down))
```

Hadoop Streaming – 使用

```python
def main():
    if sys.argv[1] == "map":
        traffic_mapper()
    elif sys.argv[1] == "reduce":
        traffic_reducer()

if __name__ == "__main__":
    main()
```

运行演示

```
mapred streaming \
  -input /user/public/week2/ \
  -output /user/public/week2_streaming/ \
  -mapper '/usr/bin/python3 traffic_stat.py map' \
  -reducer '/usr/bin/python3 traffic_stat.py reduce' \
  -file traffic_stat.py \
  -numReduceTasks 1
```

Yarn UI

https://knox.c-01b94588f59c7655.cn-hangzhou.emr.aliyuncs.com:8443/gateway/cluster-topo/yarn/cluster/apps

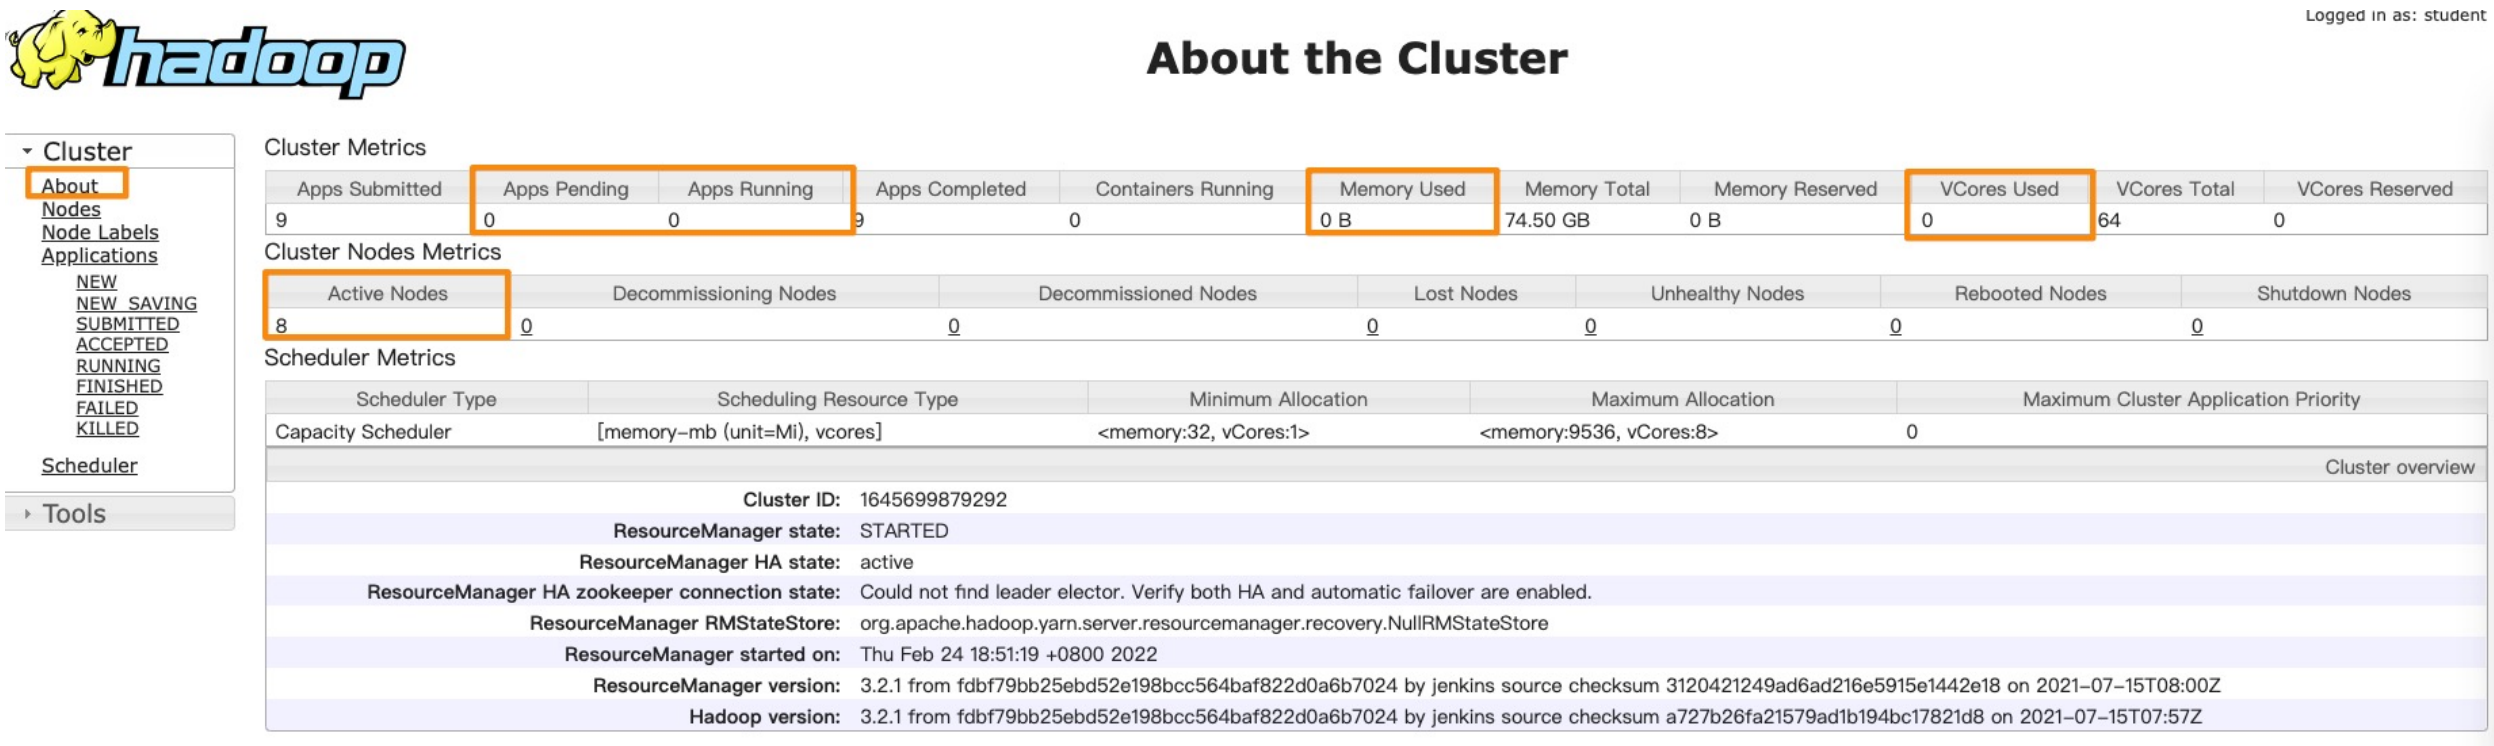

Yarn UI-Applications

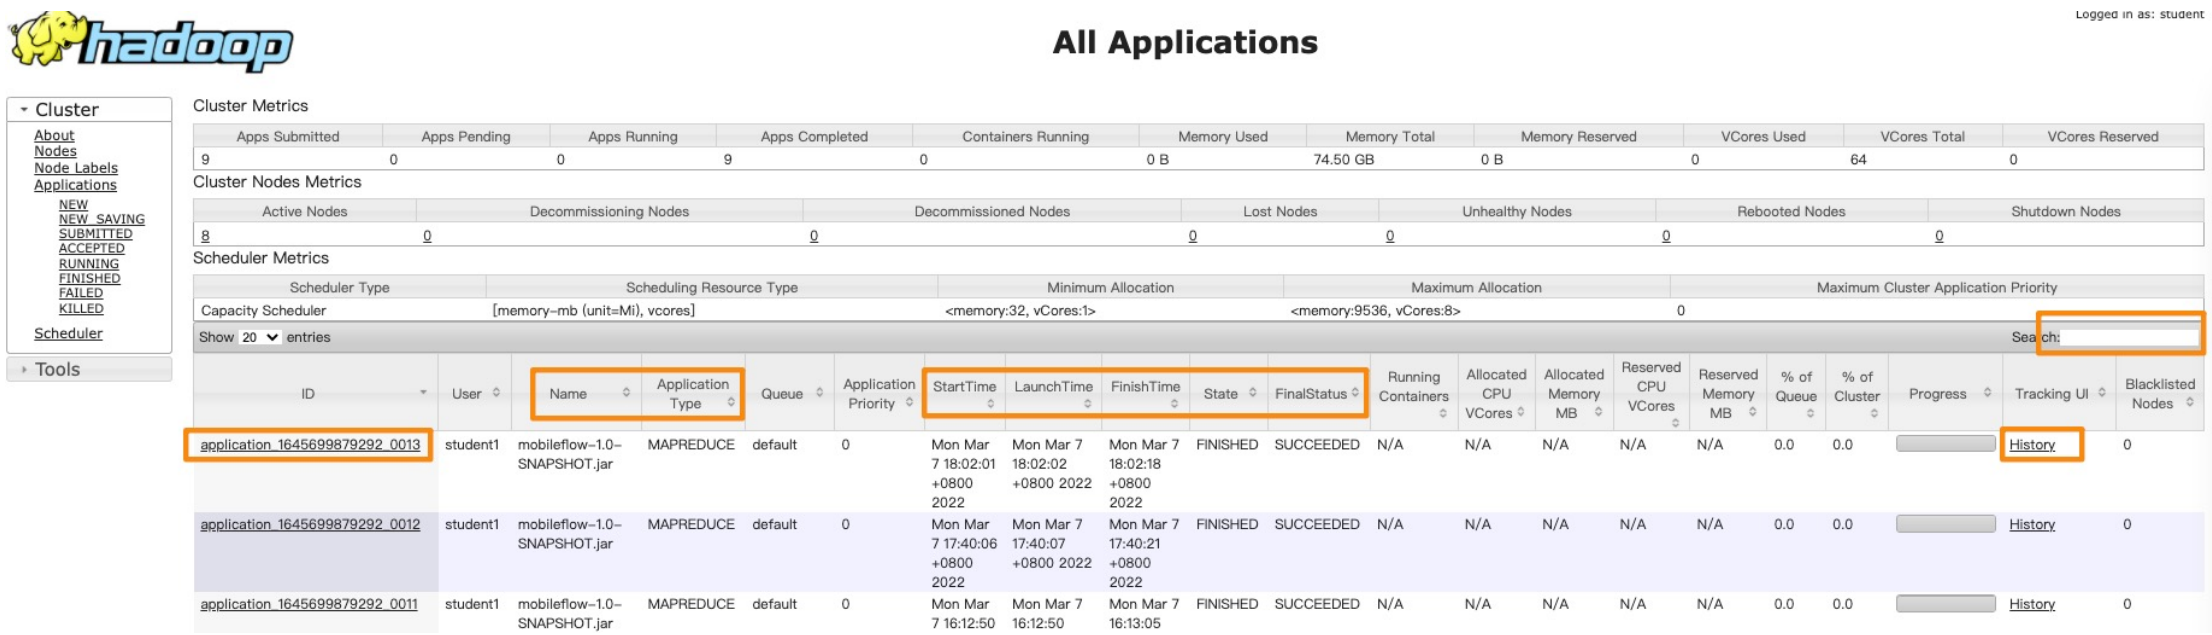

Yarn UI-Applications

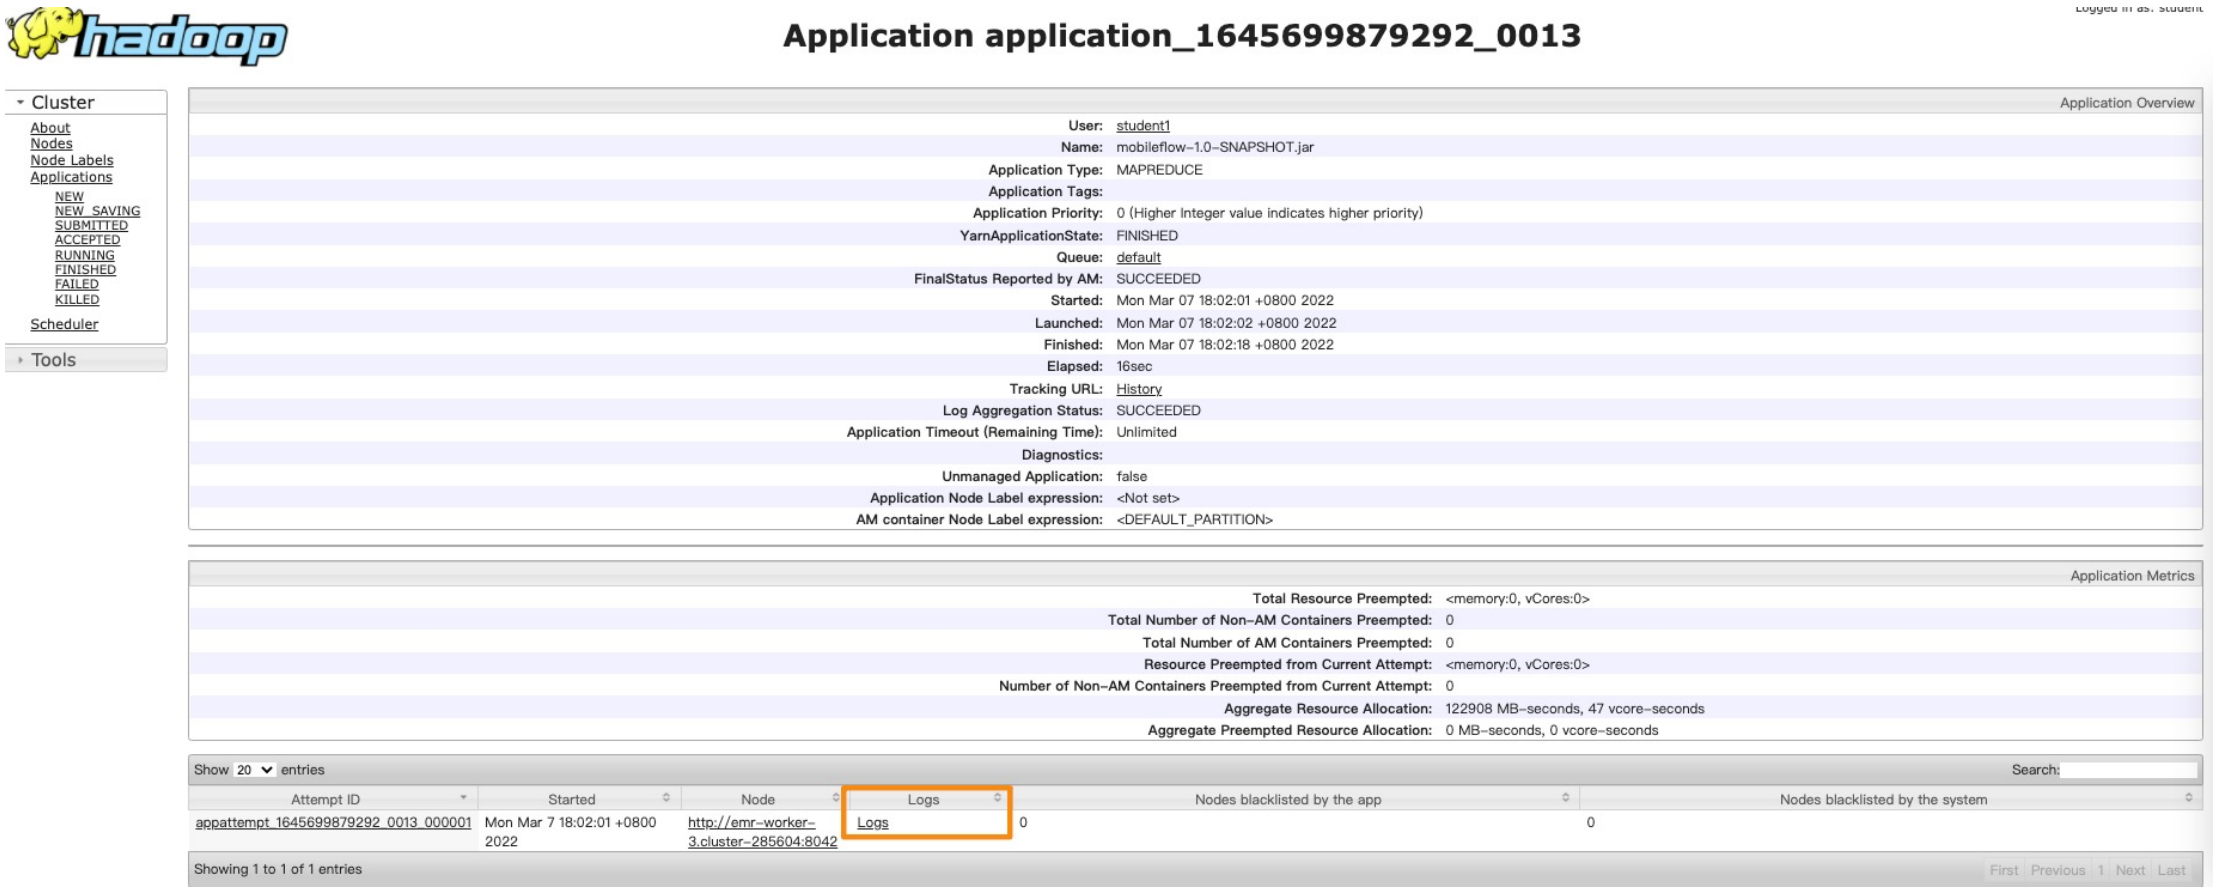

Yarn UI-Applications

Yarn UI-Applications

MapReduce Job job_1645699879292_0013

- Application
- Job
  - Overview
  - Counters
  - Configuration
  - Map tasks
  - Reduce tasks
- Tools

Job Name:	mobileflow-1.0-SNAPSHOT.jar
User Name:	student1
Queue:	default
State:	SUCCEEDED
Uberized:	false
Submitted:	Mon Mar 07 18:02:01 CST 2022
Started:	Mon Mar 07 18:02:07 CST 2022
Finished:	Mon Mar 07 18:02:18 CST 2022
Elapsed:	10sec
Diagnostics:	
Average Map Time	2sec
Average Shuffle Time	2sec
Average Merge Time	0sec
Average Reduce Time	0sec

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Mon Mar 07 18:02:03 CST 2022	emr-worker-3.cluster-285604:8042	https://knox.c-01b94588f59c7655.cn-hangzhou.emr.aliyuncs.com:8443/gateway/jobhistoryui/jobhistory/logs

Task Type	Total	Complete
Map	1	1
Reduce	7	7

Attempt Type	Failed	Killed	Successful
Maps	0	0	1
Reduces	0	0	7

Yarn UI-Scheduler

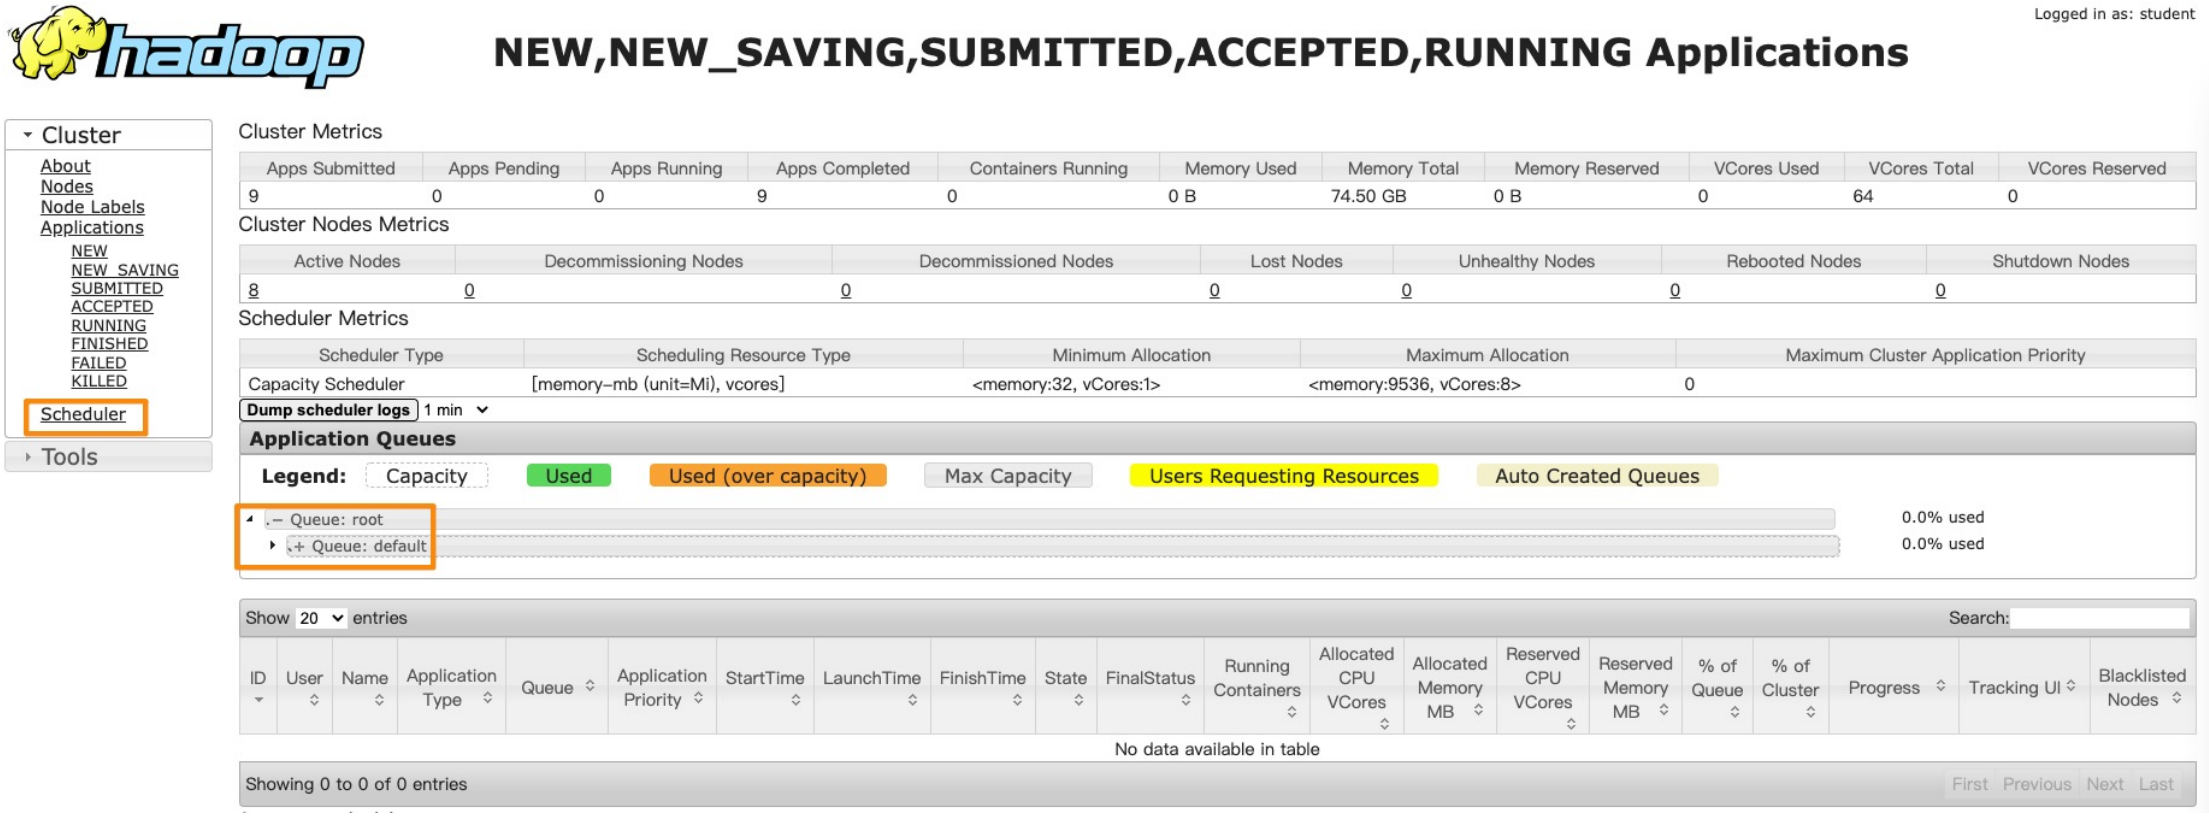

QA